

Branch: master ▾ **Big-Data** / README.md

Find file

Copy path

 jb-diplom links to images

4b2503a 1 minute ago

[1 contributor](#)

148 lines (126 sloc) 13.1 KB

Raw

Blame

History



Big-Data

Python project to scrape RSS news feeds, derive topic maps and analyze sentiment. Multiple visualizations are displayed in a Jupyter Notebook

Getting Started

The easiest approach is to clone this repository (<https://github.com/jb-diplom/Big-Data/edit/master/README.md>) and test the Notebook [RSSNewsIdentification.ipynb](#) on Jupyter.

Prerequisites

To run all the implemented features you'll need:

- Python 3.7 or newer
- Jupyter (version 6.0.3 was used for the project)
- Anaconda (version 1.9.12 was used for the project)
- Gensim (for corpora, topic modelling, word embeddings tools and downloaders)

- Natural Language Toolkit (NLTK) for text processing (Lemmatization)
- VaderSentiment (Sentiment analysis)
- Bokeh (visualization widgets and coloring)
- Feedparser (scraping of RSS Feeds)
- pyLDAvis (visualization of LDA analysis)
- ipyvolume (3d interactive scatterplots)
- ipydatawidgets (required for the interactive parts of the notebook and to ensure correct display of the progress bars)
- Beautiful Soup (parsing and cleaning of HTML/XML)
- sklearn (spectral clustering and dimension reduction using the PCA method)

Installation

Several Python packages are necessary to complete analysis/computations and for an adequate visualization in Jupyter notebook

- `conda install bokeh`
- `conda install feedparser`
- `conda install genism`
- `conda install -c conda-forge pyldavis` see also <https://github.com/bmabey/pyLDAvis>
- `conda install -c conda-forge ipyvolume`
- `conda install -c conda-forge nodejs`
- `conda install -c conda-forge ipywidgets`
- `conda install -c conda-forge ipydatawidgets`
- `import nltk`
- `nltk.download()` --> choose Models --> vader_lexicon
- `nltk.download()` --> choose Models --> punkt
- `nltk.download()` --> choose Corpora --> wordnet

- `nltk.download()` --> choose Corpora --> stopwords

Usage

The central element for accessing the functionality of the project work is by loading the [RSSNewsIdentification.ipynb](#) Jupyter Notebook. Virtually all of the useful configurative settings can be specified or modified in the run parameters, an editable Python dictionary in a cell near the beginning of the notebook. Plausible parameters are documented as comments in the code (see below).

The Run Parameters

```
runParams={'allFeeds':      getFeedDict(),    # A collection of 50 URLs for RSS Feeds, where the names can be
chosen freely                                     # example: {'Buzzfeed': 'https://www.buzzfeed.com/world.xml',
                                                    #           'Al Jazeera':
'http://www.aljazeera.com/xml/rss/all.xml'}
            'collectFeeds': False,           # Set to True to do more data-scraping
            'inputDir':     './data',         # Directory containing pre-collected RSS-Feed data in pickle
format
            'articleLimit': 10000,           # Optional specification limiting number of articles to process
(for test                                         # purposes only)
            'numAuthDispl': 30,              # specify limit of number of authors to display/plot
            'numTagsDispl': 30,              # specify limit of number of tags to display/plot
            'numTopics':    30,              # specify limit of number of topics to be found via TFIDF
Vectorization
            'numLDATopics': 30,              # specify limit of number of topics to be found via Latent
Dirichlet Allocation
            'numTopicsDispl': 30,            # specify limit of number of topics to display/plot
            'removeTopics': True,           # specify if articles now referenced by topics should be removed
(saves                                           # computation time later)
            'ngramRange':    (3,3),          # define min and max ngrams as tuple for deriving topic maps
            'useModel':      ' glove_wiki_english_F01', # Word Embedding model for soft cosine similarity
```

```

weModel : glove-wiki-gigaword-50 , # word embedding model for soft cosine similarity
calculation
subwords-300
results, load
a download of
hour
    'thresholdFuzzy': 60,
comparison) in percent
    'thresholdCosine':0.60,
(0.0 to 1.0)
    'fuzzyDocLimit': None,
relevance of
    'matrixDir':      './outdata',
    'saveMatrix':      False}
# options are e.g. glove-wiki-gigaword-50,fasttext-wiki-news-
# CAUTION: although higher dimension models produce more nuanced
# times are VERY long. In particular the first time use requires
# massive amounts of data which can easily take as much as an
# define cut-off threshold for fuzzy (Levenshtein Distance
# define cut-off threshold for soft cosine similarity as fraction
# optional integer limit to number of articles used for fuzzy
# topics (only for test purposes: default None)
# relative path for saving matrices to file
# whether to save matrix to file or not

```

A few Words on Performance

Due to the potentially large volume of data involved, running the notebook in its entirety can take many hours. This may yield very interesting results, but you may alternatively be interested in a quick overview. Some key parameters and options have been designed to help you (and especially me) to run through tests fairly quickly. The following settings are critical for performance and running times:

- `weModel` : Tests have shown - not surprisingly - far better quality and nuance of results with higher dimension word embeddings. However, start times and running times in calculating the soft cosine similarity matrix easily run into hours (or a dead end) with `fasttext-wiki-news-subwords-300` . For speed, definitely use the preconfigured `glove-wiki-gigaword-50` .
- `removeTopics` : If set to `False` , then the running time will be typically (depending also on number of topics required and specified thresholds) a factor 3 or 4 longer. Theoretically you may miss out on some possible clustering which the

soft cosine method finds but TF-IDF doesn't by the removal of articles, but it's probably sufficient to specify a few topics more to compensate and save hours of computational time, by using this flag.

- `collectFeeds` : If you just want to play with the test data and don't need fresh articles, the scraping stage can be skipped by setting this to `False`
- `articleLimit` : By setting this to 1000 or less you can get a quick overview of the analysis results. If you want to be sure to miss none of the article content, you have to set this to `None` , though. Probably time to grab a capuccino, in that case :-)
- `thresholdFuzzy` and `thresholdCosine` : in practice you get to see some optically interesting results with lower values here. However this costs time and tends to include clustering which will effectively produce more noise than anything else.

Test Data

Two crucial sets of data are supplied to help the user get up and running very quickly.

1. A set of 50 fully tested RSS-Feed URLs, which can be conveniently loaded from the function `getFeedDict()` . The user can simply specify custom values in an equivalently structured dictionary, .e.g.

```
{'Buzzfeed': 'https://www.buzzfeed.com/world.xml',  
'Al Jazeera': 'http://www.aljazeera.com/xml/rss/all.xml',  
'The Guardian': 'https://www.theguardian.com/world/rss',  
'CNBC': 'https://www.cnn.com/id/100727362/device/rss/rss.html',  
'RT': 'https://www.rt.com/rss/news/'}
```

2. A corpus of over 11000 articles which have been collected between 02.05.2020 and 31.05.2020 daily from the 50 RSS-feeds mentioned in 1.

An article limit `articleLimit` can be specified for test purposes to take a randomized sample of documents from the test corpus. It should be mentioned that with 11000 documents the runtime for the complete notebook may run into into several hours, dependant on CPU power and RAM.

Features

Initially the user can load data from a pre-configured set of 50 RSS-Feeds, or specify custom feeds for scraping. This takes typically around 75s (can be skipped using `runParam collectFeeds`). Additionally there is a large corpus of test data in the directory [data directory](#) which may be loaded for test purposes.


The steps which can be walked through are:


1. Custom configuration of `runParam`
2. Data collection from the specified RSS-feeds. The collected data is saved in pickle format to the [data directory](#)
3. All data from the test corpus in [data directory](#) are loaded. The number of articles may be reduced according to `runParam articleLimit`. The HTML content is parsed and stripped of tags, some data cleaning occurs to remove some articles which are known to not contain single stories, but rather collections of all the day's news.
4. A brief summary of the articles scraped/loaded is displayed in tabular form as a quick check (note that a more thorough view of the content is available later using the [content viewer](#))
5. To get an impression of the range of publish dates and article sizes per feed a [box-plot](#) is created
6. A statistical overview of the most prolific authors is displayed
7. Statistical overviews of the most used tags - as specified by the publishers - are displayed ([bar chart and scatterplot](#))
8. Stop words are removed from the content, TF-IDF vectorization with Lemmatization is used to convert the amassed content of the corpus into documents to a matrix of TF-IDF features deriving the specified number of topics. TF-IDF Transformer additionally scales down the impact of tokens that occur very frequently in the given corpus and that are hence empirically less informative than features occurring in a small fraction of the corpus. The user can specify not only the preferred number of features but also the n-gram range
9. The topics from 8. are compared using a Levenshtein Distance calculation to each article in the corpus. Topics within a relevant tolerance (a threshold level is specified per parameter) are held as a reference with each article including the value for relevance as a percentage. Articles may have - any often do have - references to several topics.
10. An overview of the most frequent topics and frequency of occurrence for each feed publisher are displayed in a scatter plot,
11. The overall topic frequency is visualized in a histogram

12. The soft cosine similarity matrix is derived for the corpus to ascertain similarities between every combinational pair of articles. This employs word embeddings for the calculation of the comparison vectors (see also <https://www.machinelearningplus.com/nlp/cosine-similarity/>)
13. Using the soft cosine similarity matrix, multiple [3d visualizations](#) of the clusters are produced. The colourization is simply achieved through the application of spectral clustering, a representation of the coordinates of the "bubbles" is achieved through dimension reduction using the PCA method, to create an optical separation according to the cosine similarity. The size of each bubble is calculated to be proportional to the highest topical relevance from the article's row taken from the similarity matrix. The user may specify the number of topics to display and the cut-off threshold for deciding if an article should be represented in the plot.
14. A specified number of topics are representen in the [pyLDavis visualization](#).
15. A sentiment analysis of each article is conducted using Vader. A [3d scatterplot](#) of the sentiment for topically grouped sets of articles. The coordinates are derived from the positiv, negativ and neutral sentiment values. The colour represents the compound sentiment. The size of the bubbles are exponentially proportional to the strength of topicality (cf topical relevance percentage in 9. above) for the preferred topic which can be conveniently chosen from a pull-down list of possible topics. A change of topic triggers a redraw of the 3d-visualization.
16. Finally the user can browse through the article content grouped by topic. The article content together with a subset of meta-data for the article is [displayed in a viewer](#).


Some Graphics from the Notebook

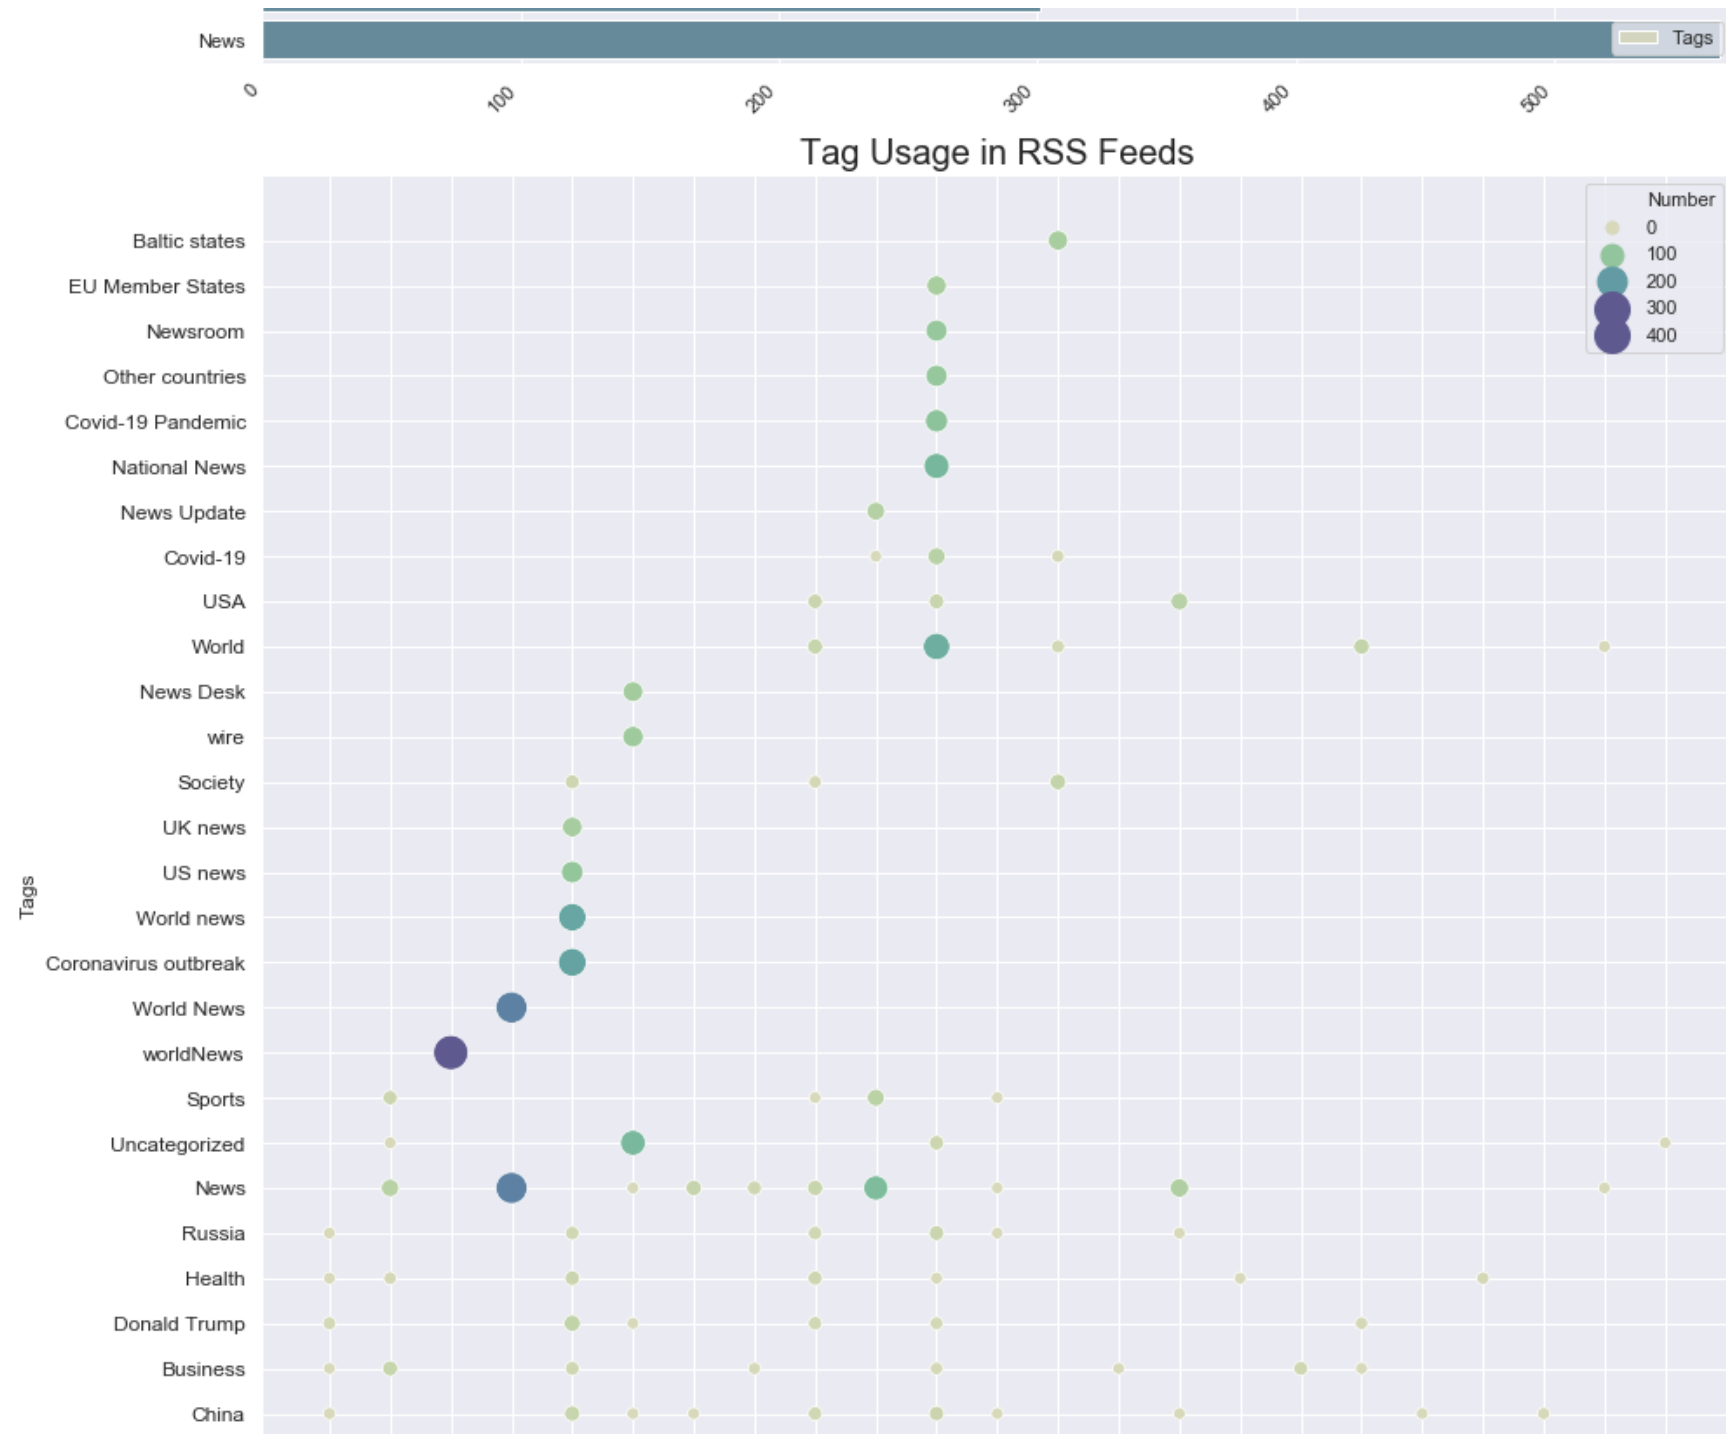
Article Sizes





Tag Usage








Cosine Similarity (3d interactive plot)

Clustering of 20 Groups, Threshold 0.45



LDA Visualization (pyLDAvis interactive viewer)

BokehJS 2.0.2 successfully loaded.


Selected Topic: Previous Topic Next Topic Clear Topic

Slide to adjust relevance metric:(2)


$\lambda = 1$

0.0 0.2 0.4 0.6 0.8 1

Intertopic Distance Map (via multidimensional scaling)




Top-30 Most Relevant Terms for Topic 2 (9% of tokens)



1. $\text{saliency}(\text{term } w) = \text{frequency}(w) * [\sum_t p(t|w) * \log(p(t|w)/p(t))]$ for topics t ; see Chuang et. al (2012)
2. $\text{relevance}(\text{term } w | \text{topic } t) = \lambda * p(w|t) + (1 - \lambda) * p(w|t)/p(w)$; see Sievert & Shirley (2014)

Sentiment Analysis (3d interactive plot)

Sentiment Analysis for Topic 'first onworld weekly'



Content Viewer

```
%aimport sentiment3d
importlib.reload(sentiment3d)
from sentiment3d import conductSentimentAnalysis, contentsViewer
contentsViewer(allDict, topics)
```

Mapping Topics: 100%  1878/1878 [00:00<00:00, 8948.45it/s]

Pre-select a topic and choose an article to view its contents

Topics: RSS Entries:

Content:

Title:	China's missing 'Bat Woman' with secrets...
Feed:	Times of India
Published:	02/05/2020, 20:03:49
Sentiment:	pos: 0.0 neg: 0.438 neut: 0.562 comp: -0.7783

China's missing 'Bat Woman' with secrets of Covid-19 origin denies defecting to West

Versioning

All code is versioned in Github in [this repository](#).

Authors

- Janice Butler - *Initial work* - [Big-Data Project](#)

License

This project is licensed under the MIT License - see the [LICENSE.md](#) file for details

Acknowledgments

- Many thanks to Damian Trilling for the ideas on using cosine similarity and Levenshtein Distance